

From Open Hardware Design to Silicon: Lessons From Some Big (Academic) Chips

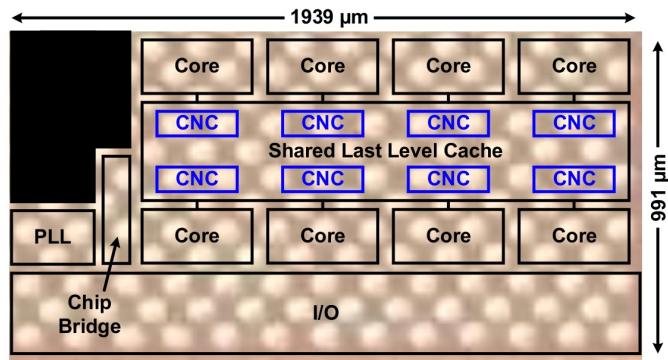
Jonathan Balkind

ArchLab, UC Santa Barbara

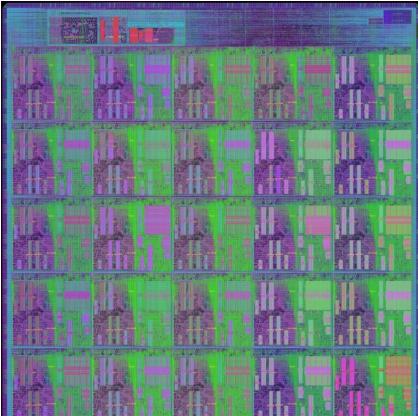
Building Realistic Scalable Prototypes

- Seen as expensive and risky
- Need open IP to reuse
- Scalability and configurability are challenging

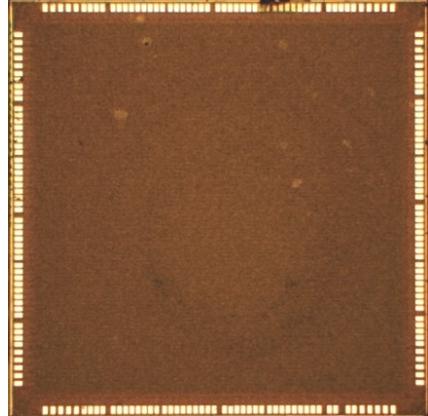
- OpenPiton tackles these issues
- For our users, OpenPiton acts as:
 - A research platform for manycore SoC development
 - A nexus for open source hardware



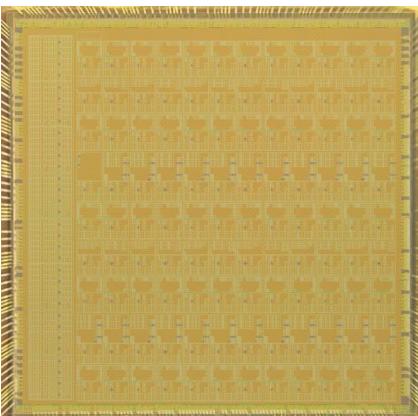
G.K.Chen et al. JSSC 2023



Piton



CIFER



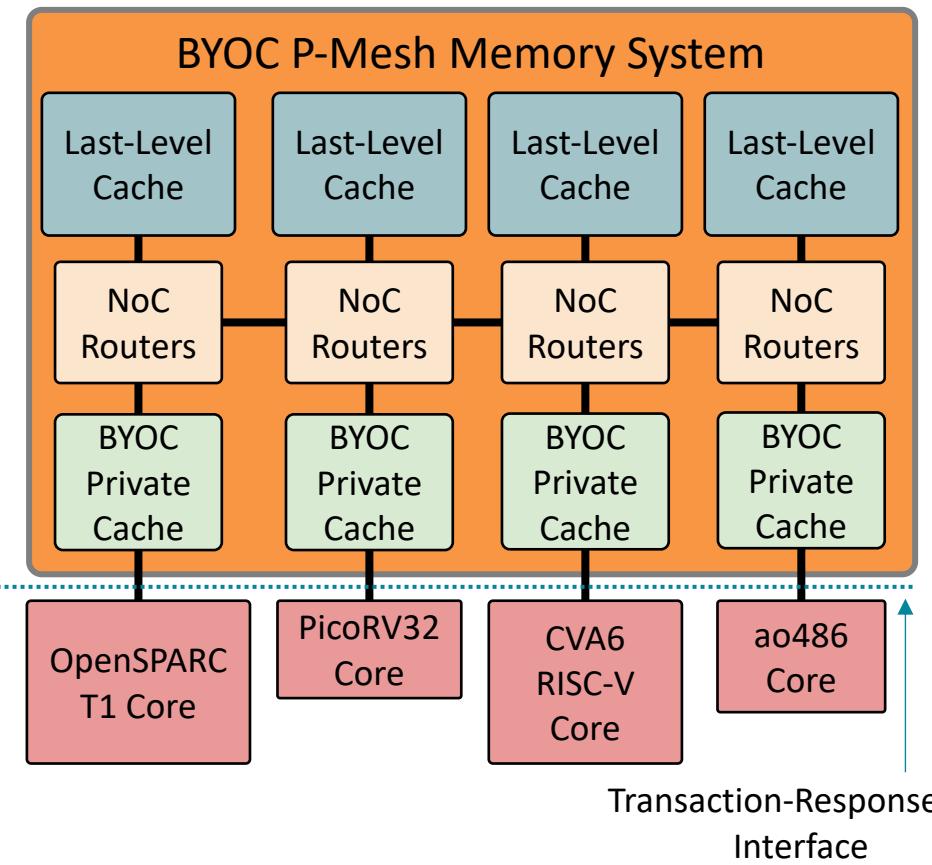
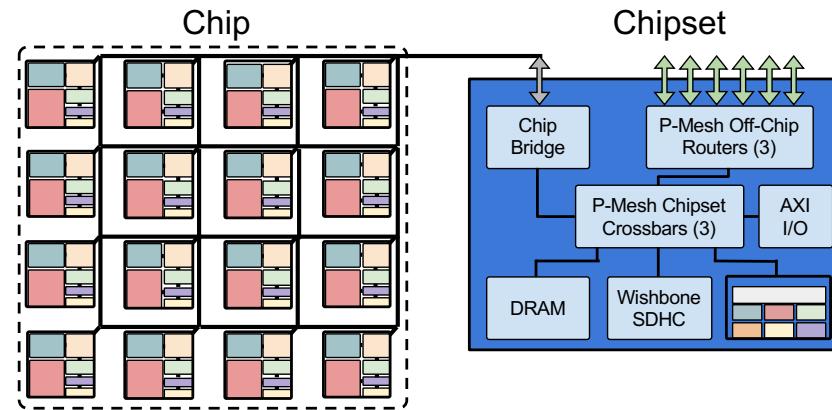
DECADES

Coming
soon!
Currently
in fab

Polara

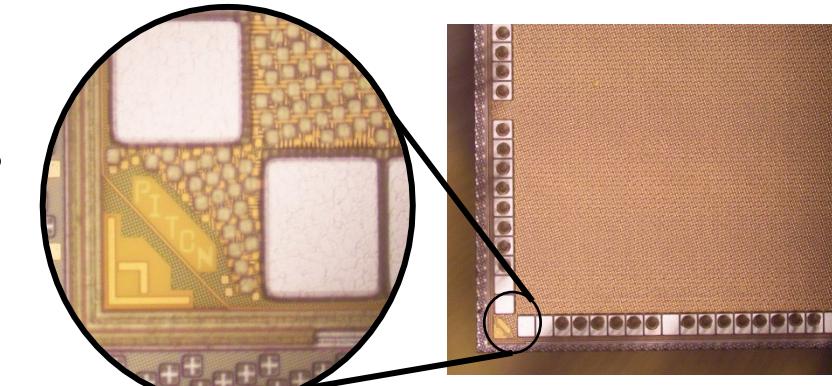
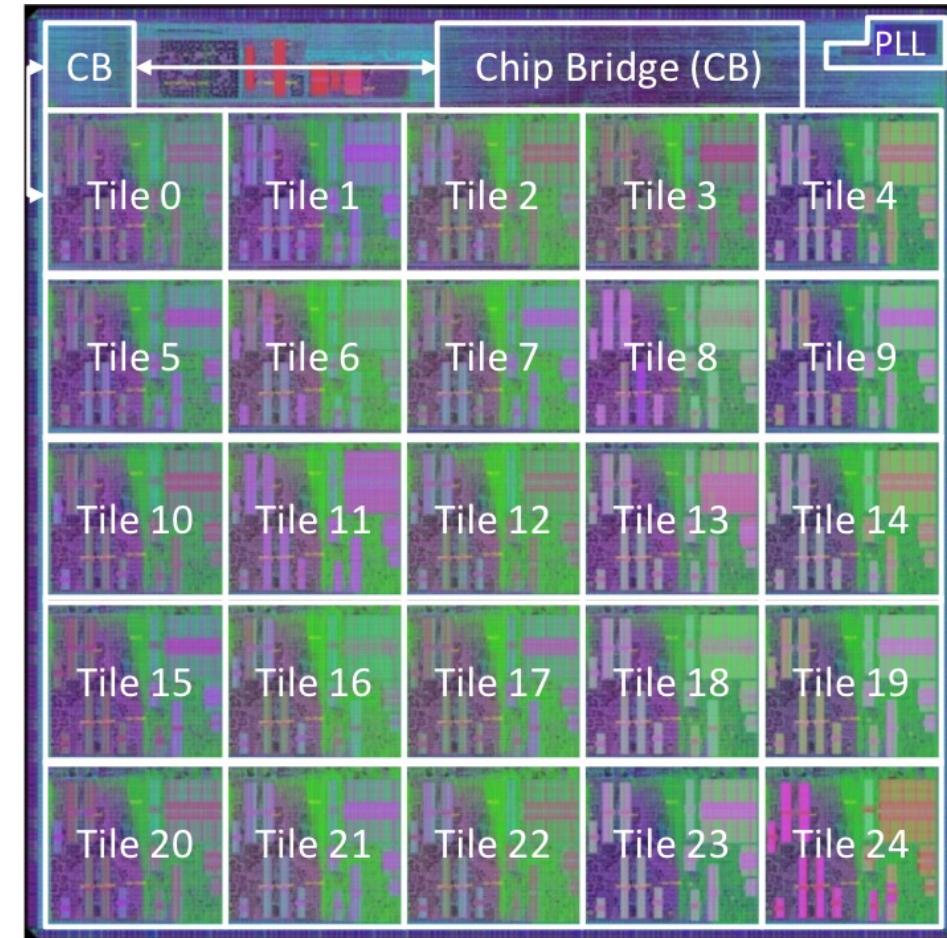
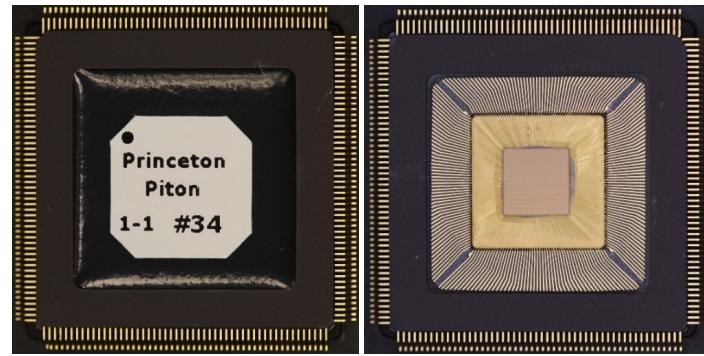
OpenPiton & P-Mesh Coherence

- Open source manycore (since 2015)
- P-Mesh coherence scales to $\frac{1}{2}$ billion cores
 - Innovation: we connect cores of different ISAs in one open-source system!
 - Core interface to cache coherence (TRI)
 - **We have connected >15 cores of six ISAs**
- Configurable core, uncore
- Simulation in VCS, ModelSim, Incisive, Verilator, Icarus, Riviera*, Vivado
- Runs full stack multi-user Debian Linux
- Used in >75 public research works

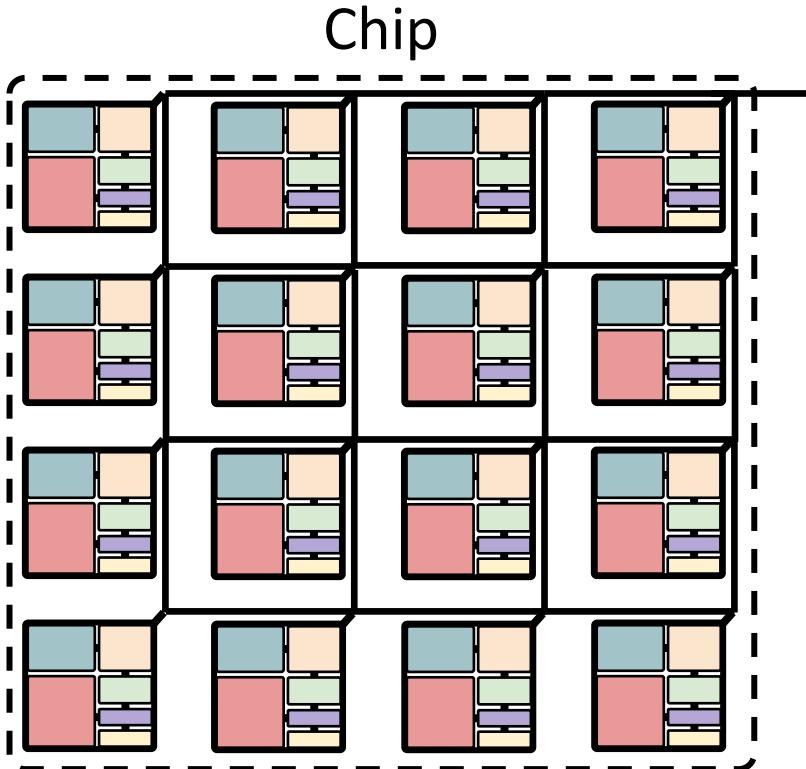


Piton Chip

- 25 cores
 - Modified 64 bit OpenSPARC T1 Core
- 3 P-Mesh NoCs
 - 64 bit, 2D Mesh
 - Extend off-chip enabling multichip systems
- P-Mesh Directory-Based Cache System
 - 64kB L2 Cache per core (Shared)
 - 8kB L1.5 & L1 Data Caches
 - 16kB L1 Instruction Cache
- IBM 32nm SOI Process
 - 6mm x 6mm
 - 460 Million Transistors - Among the largest academic chips
- Target: 1 GHz Clock @ 900 mV
- Received silicon and runs full-stack Debian in lab



System Overview - Chip



- Extending NoCs off-chip enables multi-chip
- Off-chip chipset logic was developed post-tapeout
- Packets can be arbitrarily modified off-chip
- **Lesson:** This enables agile prototyping and addition of peripherals
- **Lesson:** Over-provisioned packet format enabled research modifications

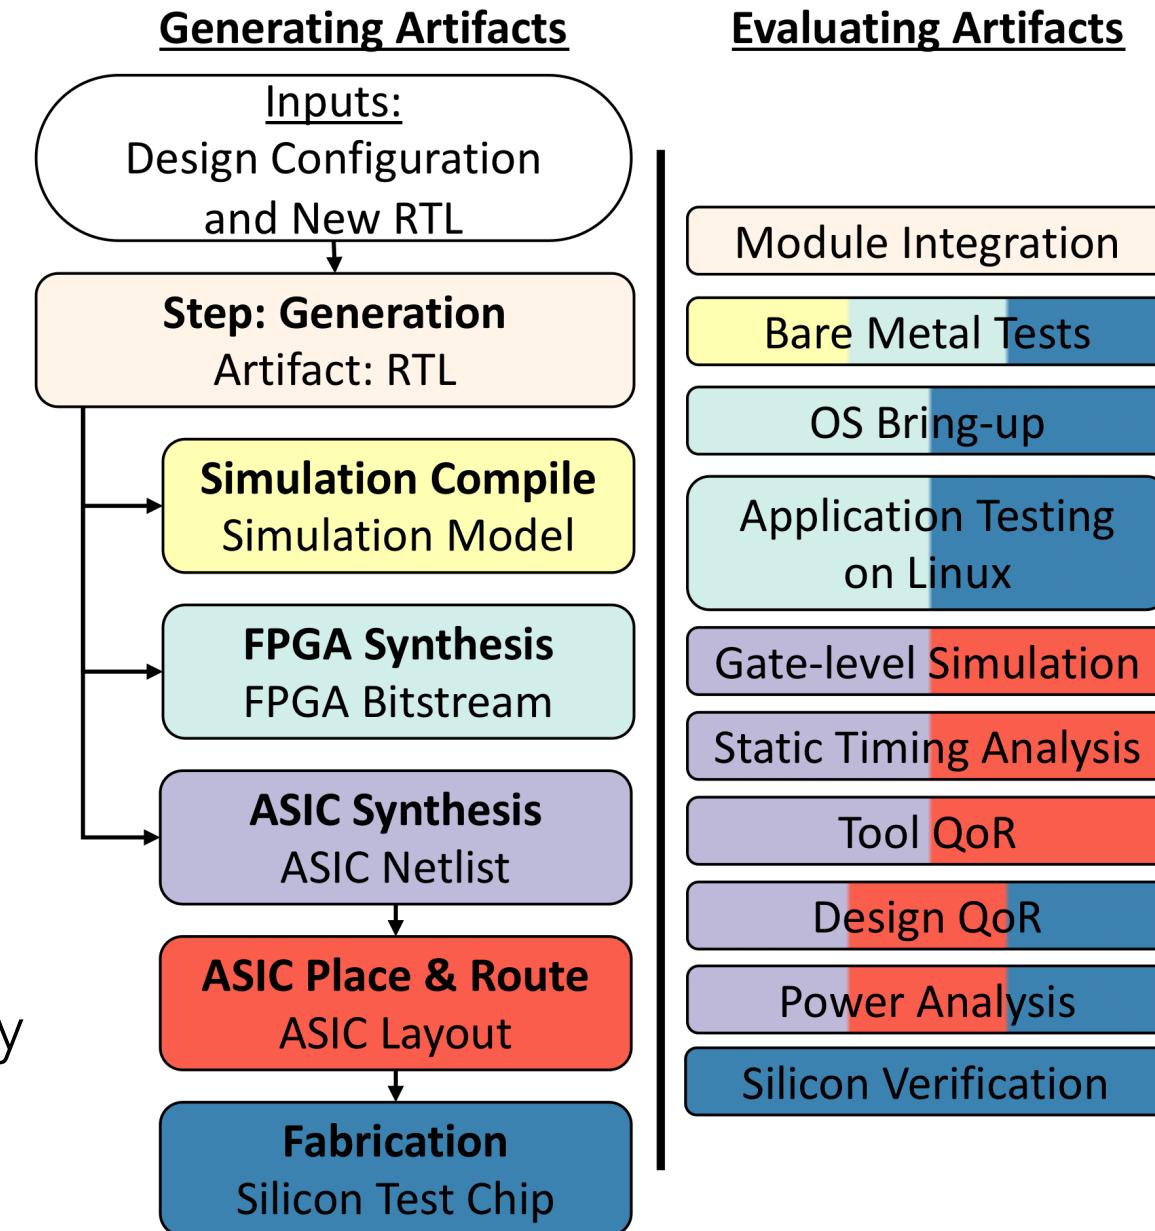
Configuration Options

- Configuration provided both in core and uncore
- Many options needed to enable design space exploration
- Parameterisation enabled by PyHP preprocessing
 - Python can be embedded in Verilog to produce complex Verilog code
- **Lesson:** PyHP and heterogeneity can be at odds
 - Leaning more on native Verilog today

Component	Configurability Options	
Cores (per chip)	Up to 65,536	
Cores (per system)	Up to 500 million	
Core Type	OpenSPARC T1	Ariane 64 bit RISC-V
Threads per Core	1/2/4	1
Floating-Point Unit	FP64/32	FP64/32/16/FP8, BFLOAT16
TLBs	8/16/32/64 entries	Number of entries
L1 I-Cache	Number of Sets, Ways	
L1 D-Cache	Number of Sets, Ways	
L1.5 Cache	Number of Sets, Ways	
L2 Cache	Number of Sets, Ways	
Intra-chip Topologies	2D Mesh, Crossbar	
Inter-chip Topologies	2D Mesh, 3D Mesh, Crossbar, Butterfly Network	
Bootloading	SD/SDHC Card, UART, RISC-V JTAG Debug	

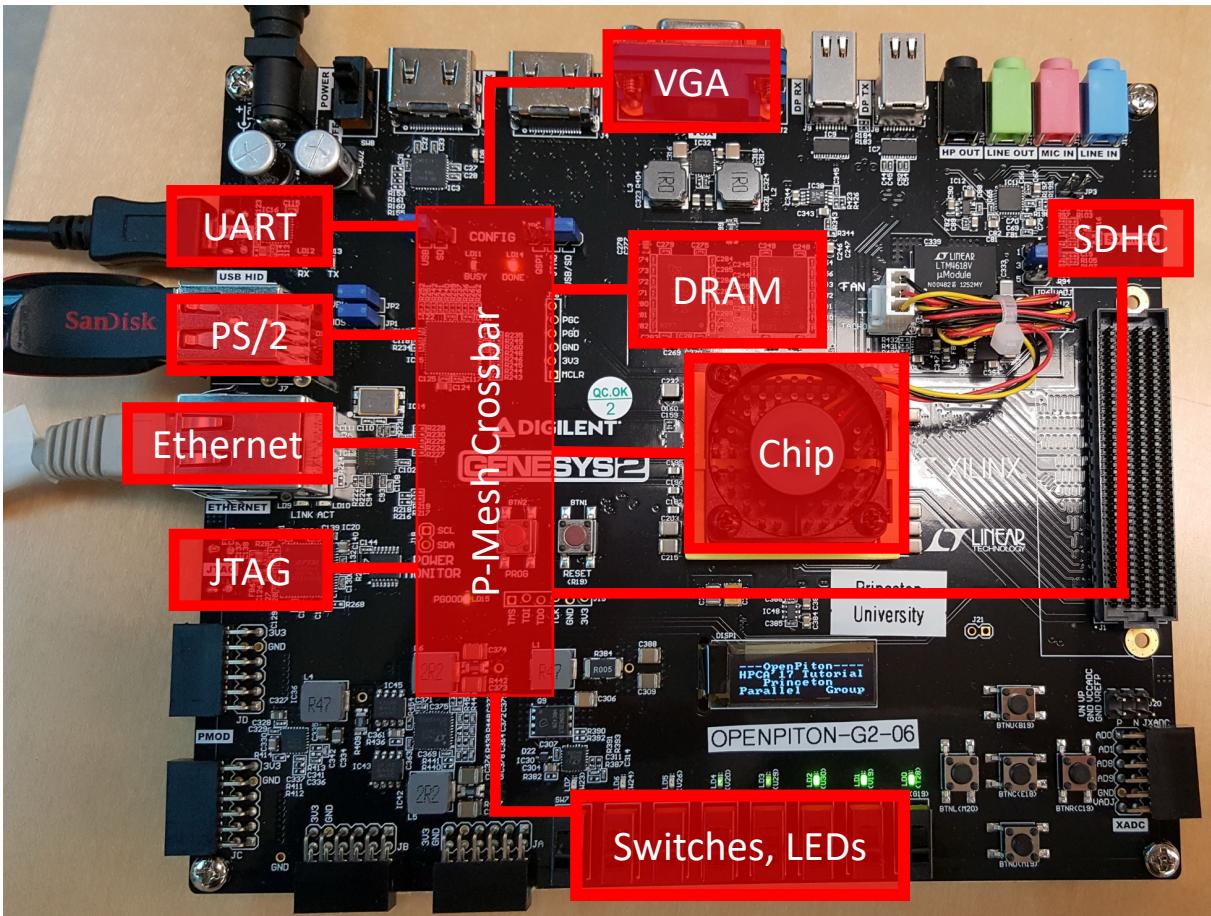
Tool Flows

- Simulation (**sims**)
 - Flow came with VCS simulation
 - Today further supports Cadence NCSim, Mentor ModelSim, Icarus Verilog, Verilator, Xilinx Vivado, simulation on MacOS
- FPGA (**protosyn**)
- Synthesis and Backend (**rftf**)
- **Lesson:** supporting all users means supporting open tools
- **Lesson:** push-button, single command flows are crucial to research productivity
- **Lesson:** porting FPGAs/ASIC libraries must be easy



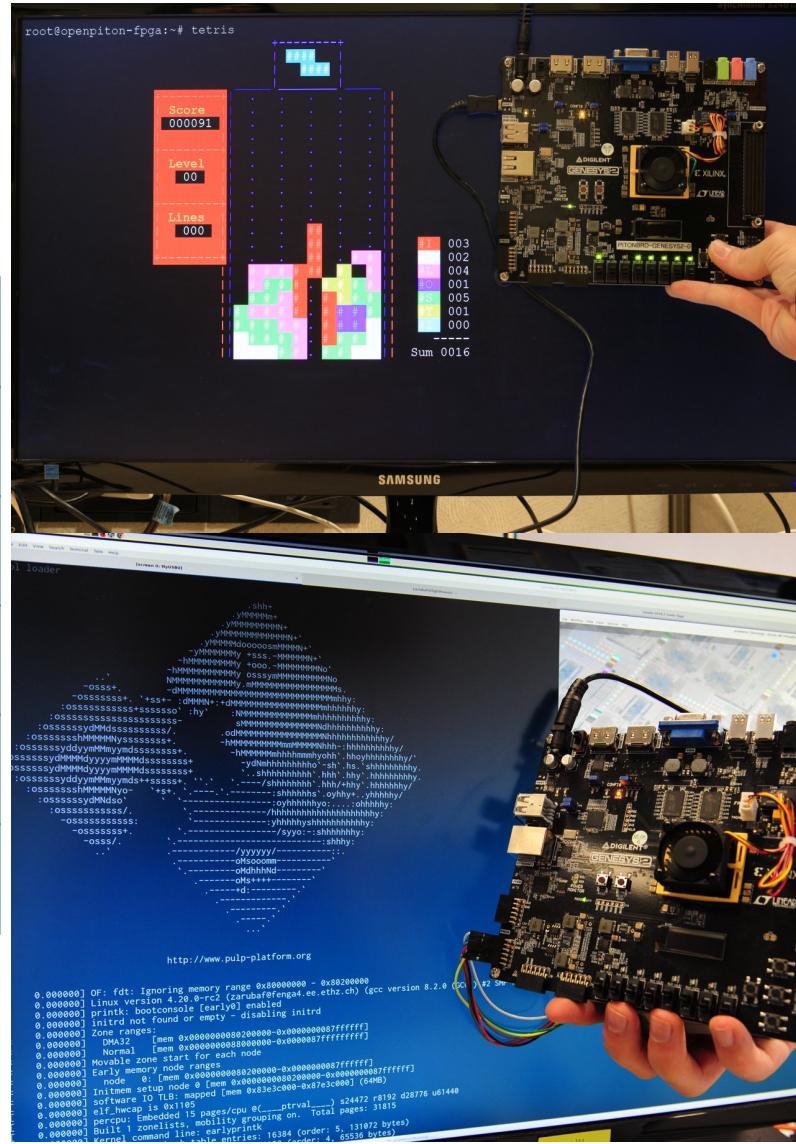
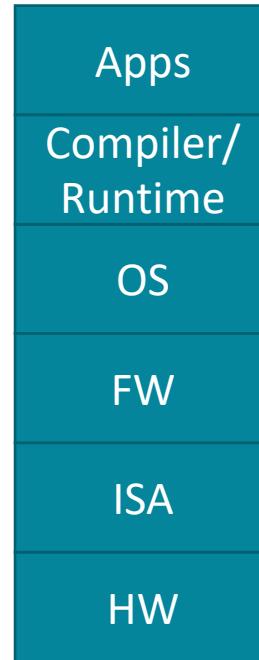
FPGA Prototyping

- Chipset crossbars generated from XML description of devices
 - Chip has no knowledge of chipset
 - Packets from chip rewritten to route to correct port
- Many devices supported for use bare metal and from Linux
- **Lesson:** Productivity enhancements for repeated menial tasks are worthwhile at scale
 - Debian boot time reduced from ~45 minutes to ~5 minutes



Open Source System Stack

- All levels of the stack are open source
- Unmodified applications can run in Debian Linux on OpenPiton
- Or users can modify at any level of the stack!
- Debian boots on SPARC + RISC-V
 - Minor Linux patches used (and open)
- Modifications to reset code and hypervisor made for SPARC



Open Source Design and Data

The diagram illustrates the OpenPiton ecosystem. At the center is a large blue cloud containing the text "OpenPiton" and the website "http://www.openpiton.org". Above the cloud, the word "Design" is written above "OpenPiton" and "Characterization Data" is written below it. To the left of the cloud is a diagram of a chip architecture with a "Chip Bridge" at the top, connected to "Tile 0", "Tile 1", "Tile 5", "Tile 10", "Tile 15", "Tile 16", "Tile 20", "Tile 21", "Tile 22", "Tile 23", and "Tile 24". A double-headed arrow connects the "Design" and "Characterization Data" sections. To the right of the cloud are three graphs. The top graph shows "Measured Power (W)" vs "Hops" for various benchmarks: "nop", "and", "and", "mix", "mix", "sqrt", "sqrt", "fadd", "fadd", "fdiv", "fdiv", "ldt", "ldt", "stx (N)", "stx (N)", and "time (T)". It includes trendlines for FSW (-16.68 pJ/hop), FSVA (-16.98 pJ/hop), HSW (-11.18 pJ/hop), and NSW (-3.58 pJ/hop). The middle graph is a bar chart of "Measured Power (W)" for SRAM Dynamic Power, Core Dynamic Power, SRAM Static Power, and Core Static Power across different memory configurations. The bottom graph is a scatter plot of "Measured Power (W)" vs "Hops" for the same benchmarks as the top graph, with trendlines for FSW, FSVA, HSW, and NSW.

Design
Characterization Data

OpenPiton

OpenPiton
<http://www.openpiton.org>

Measured Power (W)

Trendlines: FSW (-16.68 pJ/hop), FSVA (-16.98 pJ/hop), HSW (-11.18 pJ/hop), NSW (-3.58 pJ/hop)

Measured Power (W)

SRAM Dynamic Power, Core Dynamic Power, SRAM Static Power, Core Static Power

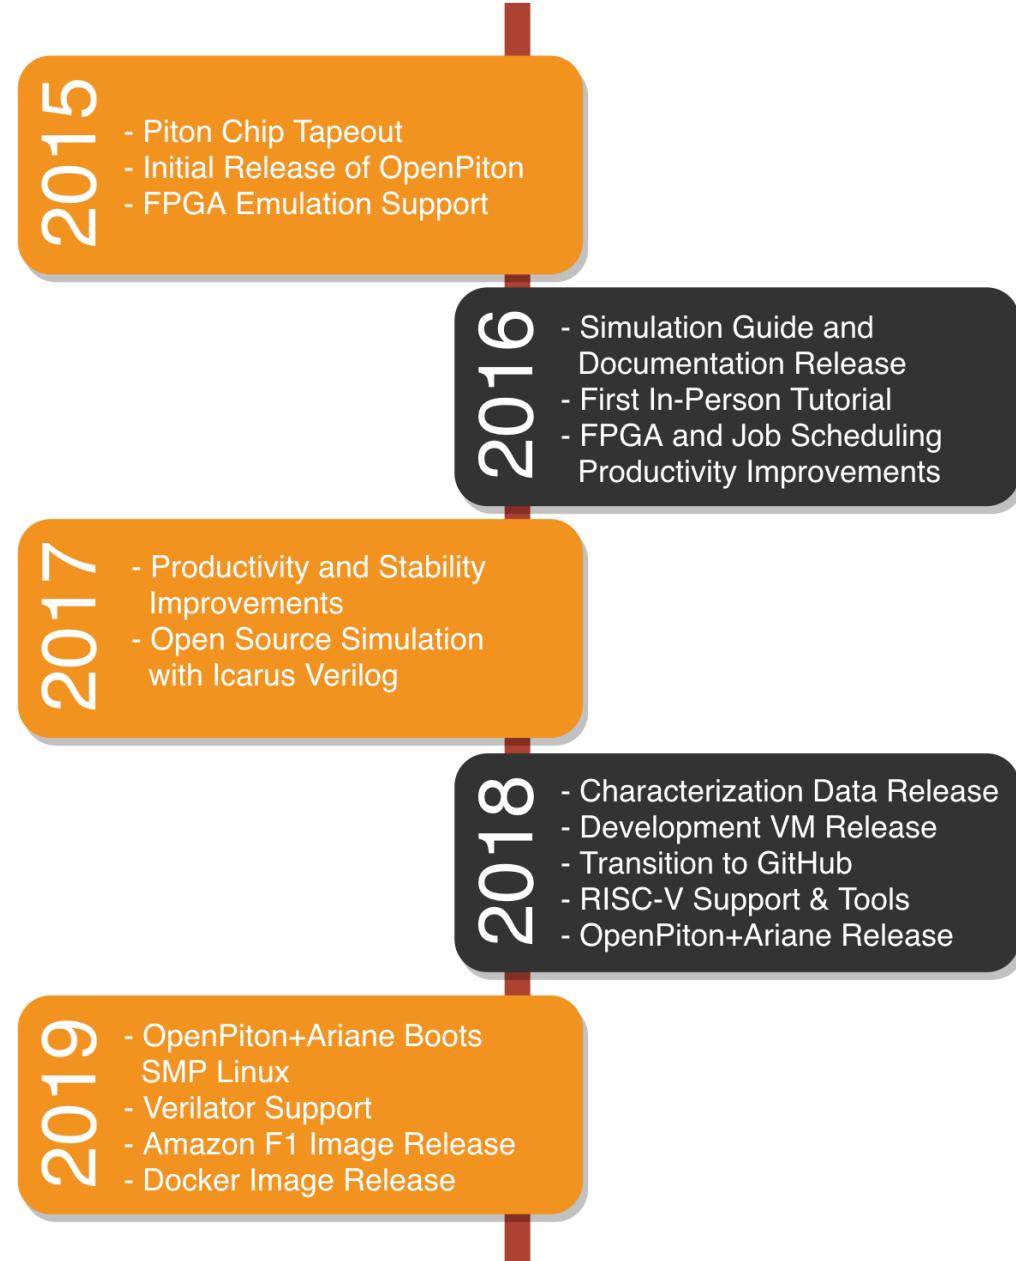
Measured Power (W)

Hops

First detailed power and energy characterization of an open source manycore!

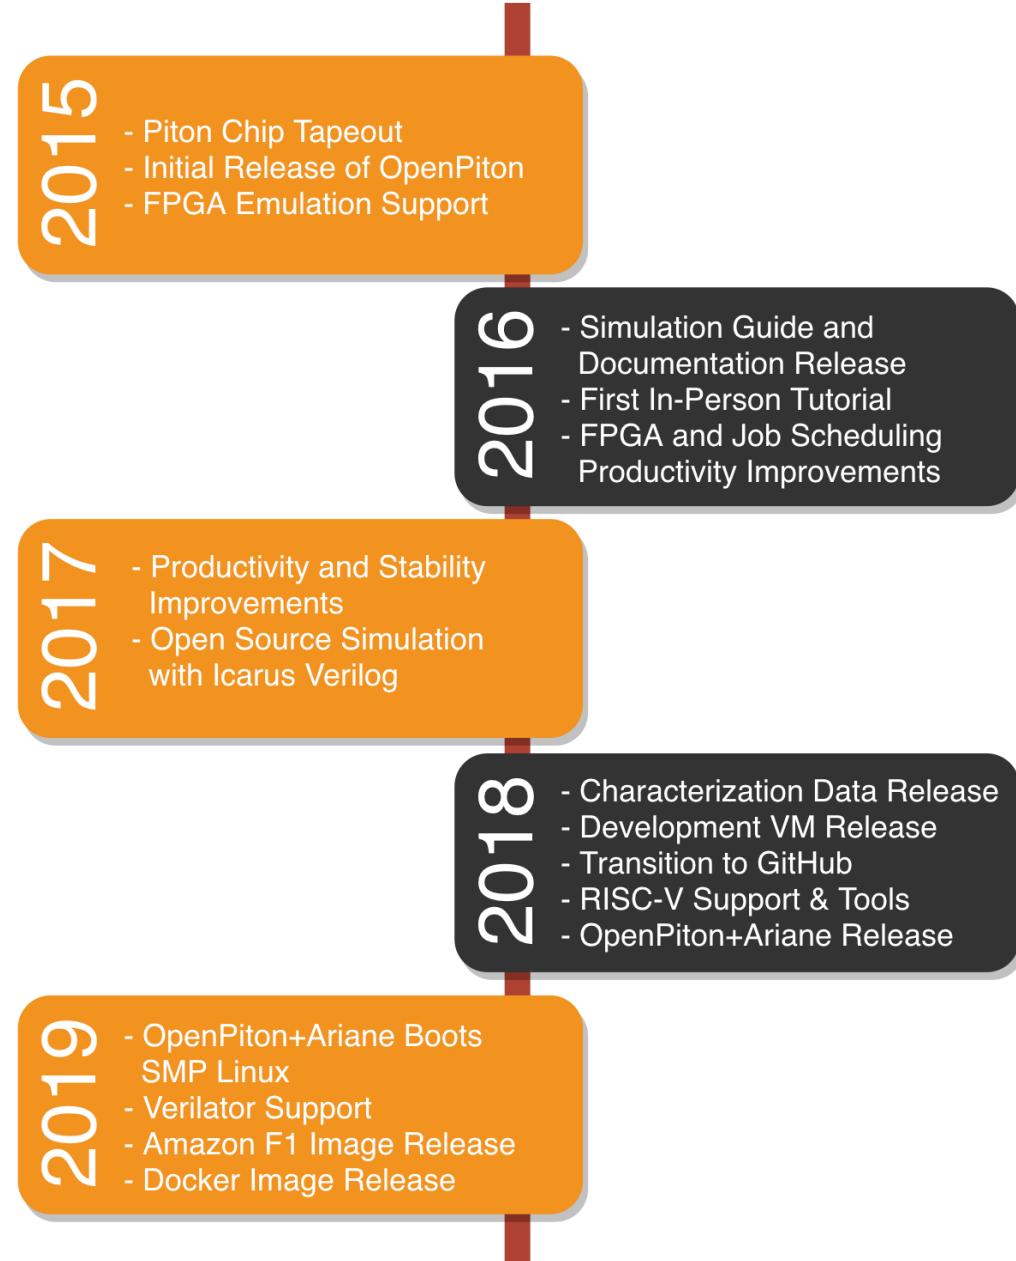
Open Source & Community

- 13 releases since 2015
- Transitioned to GitHub in 2018
- Open source ecosystem
 - Research: OpenPiton+Ariane, PyOCN, ILA
 - Tools: Icarus, Verilator, FuseSoC, Yosys, OpenTimer, OpenSTA
 - Peripherals: SDHC, DRAM, VGA, AXI
 - Industrial use: Emulation, training
- Community building
 - Thirteen tutorials
 - Mailing list for announcements
 - Google group for support



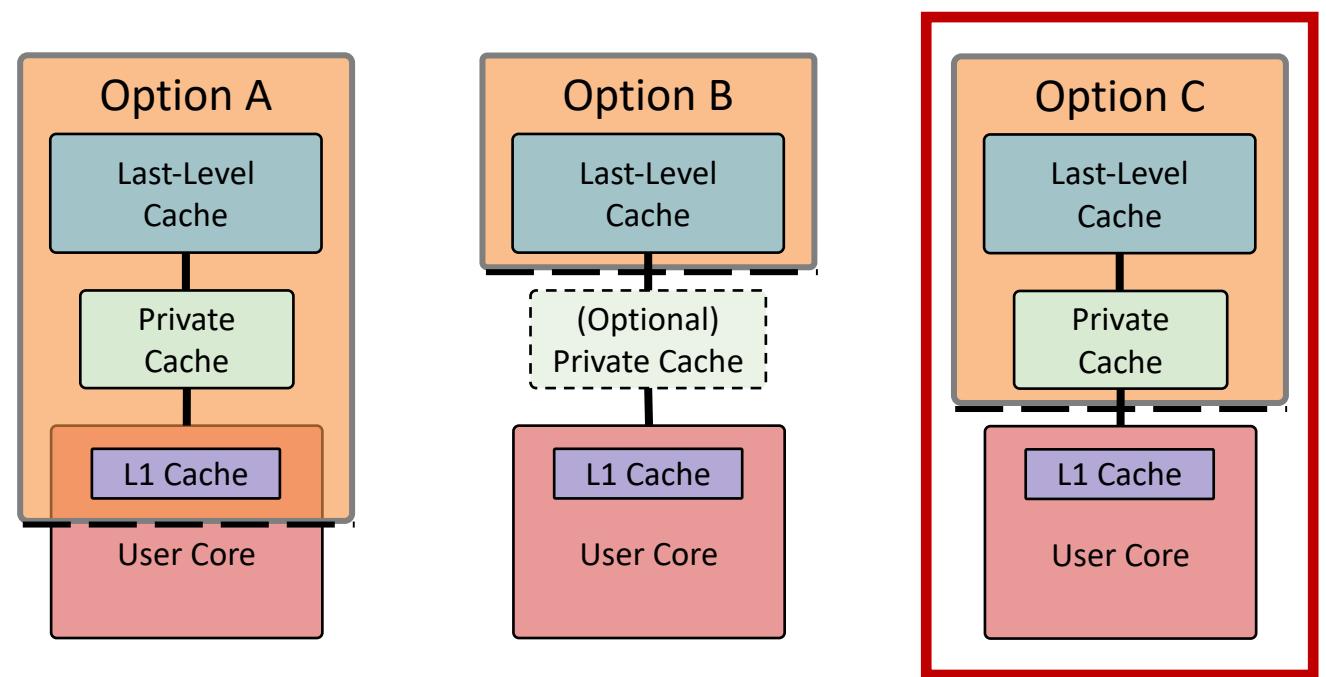
Open Source & Community

- **Lesson:** Lowering barrier to entry keeps users engaged
 - Wrote getting started guide
 - Prebuilt docker, virtual machine images
- **Lesson:** Tutorial materials increase access
 - Slides and pre-recorded tutorial available
- **Lesson:** Archiving issues helps in future
 - Users find answers and see active community
- **Lesson:** Open development improves code quality and adoption



Core Interface Options

- Several interface options with different trade-offs
- We want to minimise effort to connect new cores
- We define the Transaction-Response Interface (TRI)
 - Simple, extensible, low overhead



Provide 3 cache levels
Interface before L1

Provide 1 cache level
Interface before LLC

Provide 2 cache levels
Interface after L1

Core to Cache	Cache to Core	Subset
Load	Load Response	Load/Store
Store	Store Ack	Load/Store
Instruction Miss	Instruction Return	Instruction
	Invalidation Request	Invalidation
Atomic	Atomic Response	Atomic
Outbound Interrupt	Inbound Interrupt	Interrupt

Cores Connected With TRI Subsets

- CVA6/Ariane (RISC-V 64)
 - All subsets
- OpenSPARC T1 (SPARC V9)
 - All subsets
- BlackParrot (RISC-V 64)
 - All subsets
- PicoRV32 (RISC-V 32)
 - Load/store + Atomic
- ao486 (x86) 
 - Instruction + Load/store
- QEMU (Many ISAs)
 - Load/store
- VexRiscv (RISC-V 32)
 - Instruction
- Microwatt (OpenPOWER)
 - Instruction + Load/store*

Core to Cache	Cache to Core	Subset
Load	Load Response	Load/Store
Store	Store Ack	Load/Store
Instruction Miss	Instruction Return	Instruction
	Invalidation Request	Invalidation
Atomic	Atomic Response	Atomic
Outbound Interrupt	Inbound Interrupt	Interrupt

TRIcky Design Details

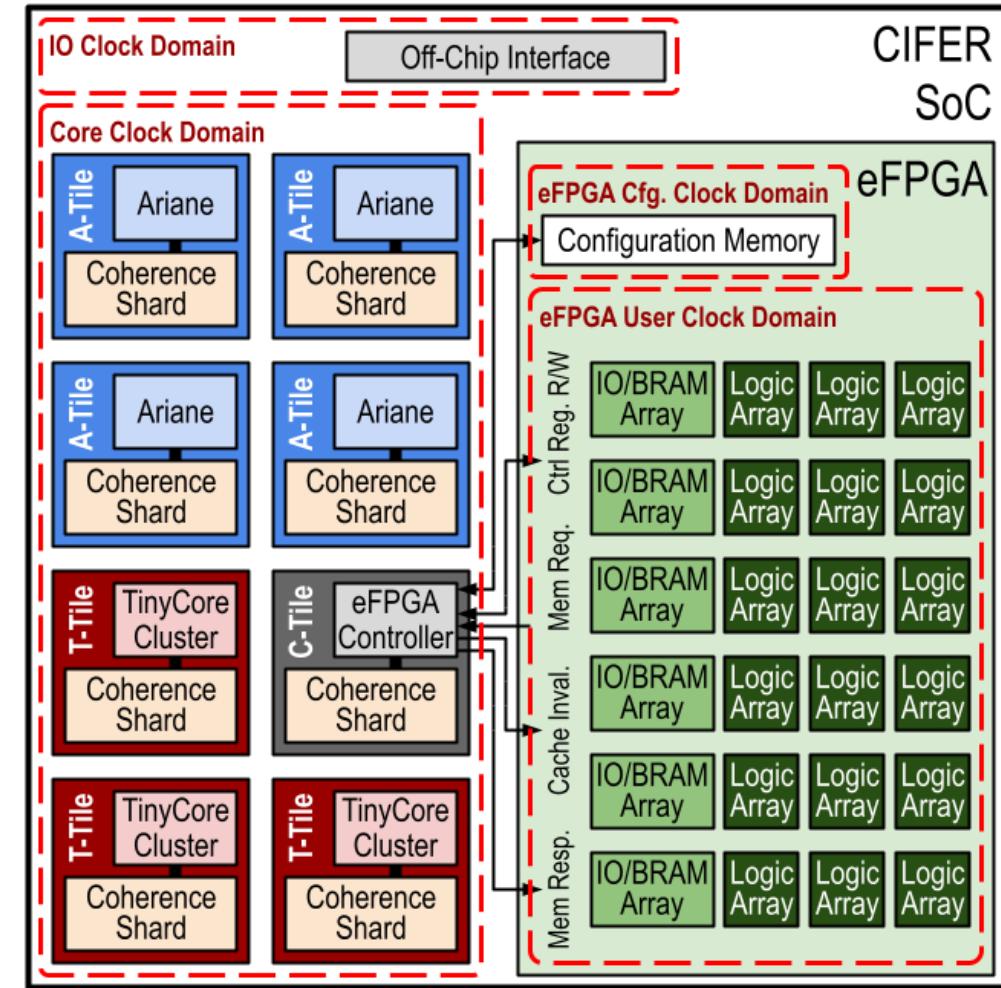
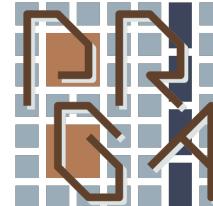
- TRI subsets
- L1 cache organisation
- Supporting existing protocols
- Configurability
- Write-back L1 caches
- Handshakes
- Reset strategy
- ISA complexity

Enabled Research

- OpenPiton team:
 - Execution Drafting [McKeown MICRO 2014]
 - Coherence Domain Restriction [Fu MICRO 2015]
 - MITTS [Zhou ISCA 2016]
 - PiCL [Nguyen MICRO 2018]
 - Piton Characterisation [McKeown HPCA 2018]
 - JuxtaPiton [Lim FPGA 2019]
 - BYOC [Balkind ASPLOS 2020]
 - MAPLE [Orenes Vera ISCA 2022]
 - Cohort [Wei ASPLOS 2023]
 - SMAPPIC [Chirkov ASPLOS 2023]
- Collaborators:
 - Oblivious RAM [Ren TDSC 2017]
 - POSH/IDEA/SDH programs
- External:
 - Hardware trojan detection [Elnaggar ITC 2017, Basu DAC 2019, Elnaggar T-VLSI 2019]
 - IoT reliability lifetime [Lerner ISQED 2017, Lerner T-VLSI 2018]
 - NCFET characterization [Rapp DAC 2019, Salamin ISLPED 2019]
 - Scan test power analysis [Indino 2017]
 - Modeling and control of EMI [Gorman 2018, Gorman MICRO 2019]
 - 3D ICs [Pentapati IEEE Micro 2019, Bamberg DATE 2020]
- Student research projects:
 - Transactional Memory, Heterogeneous ISA, MIAOW GPU, ...
- Industrial:
 - Intel, Western Digital, Google, ...

CIFER Architecture Overview

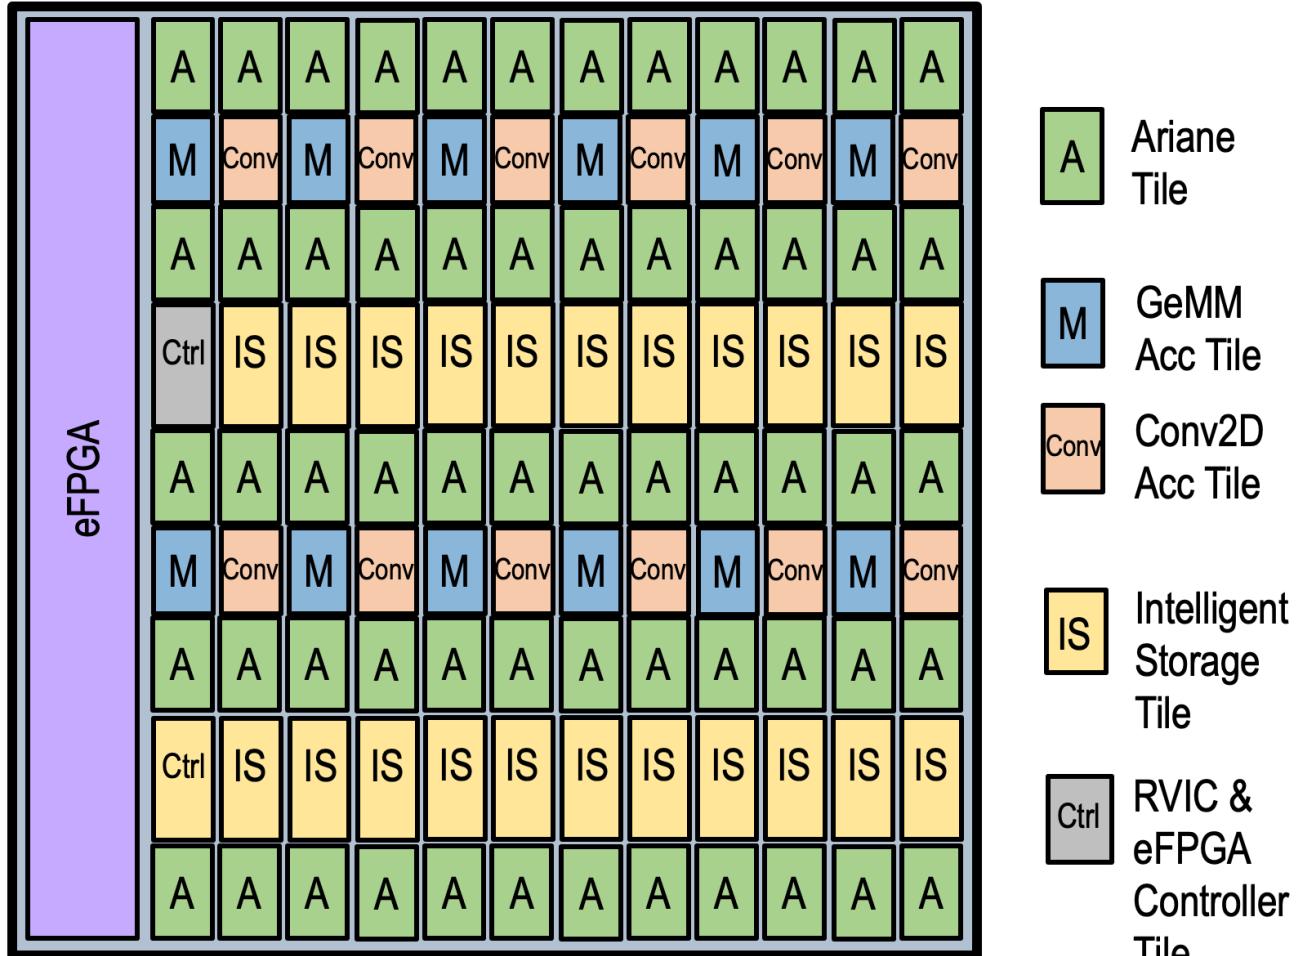
- 4×  CVA6
- 3×  Tiny-core clusters
 - 6-core, MIMD, 32-bit RISC-V (18 in total)
- Embedded FPGA
 - RTL design with PRGA
 - Synthesized with standard cells
 - 6,720 basic logic elements (LUT6 + 2× DFF)
 - 54KB BRAM
- 12nm FinFET, 456M transistors, 4x4mm
 - V_{nom} [$V_{min} - V_{max}$] (V): 0.80 [0.68 – 1.10]
 - $F_{max}@V_{max}$ (GHz): 1.2



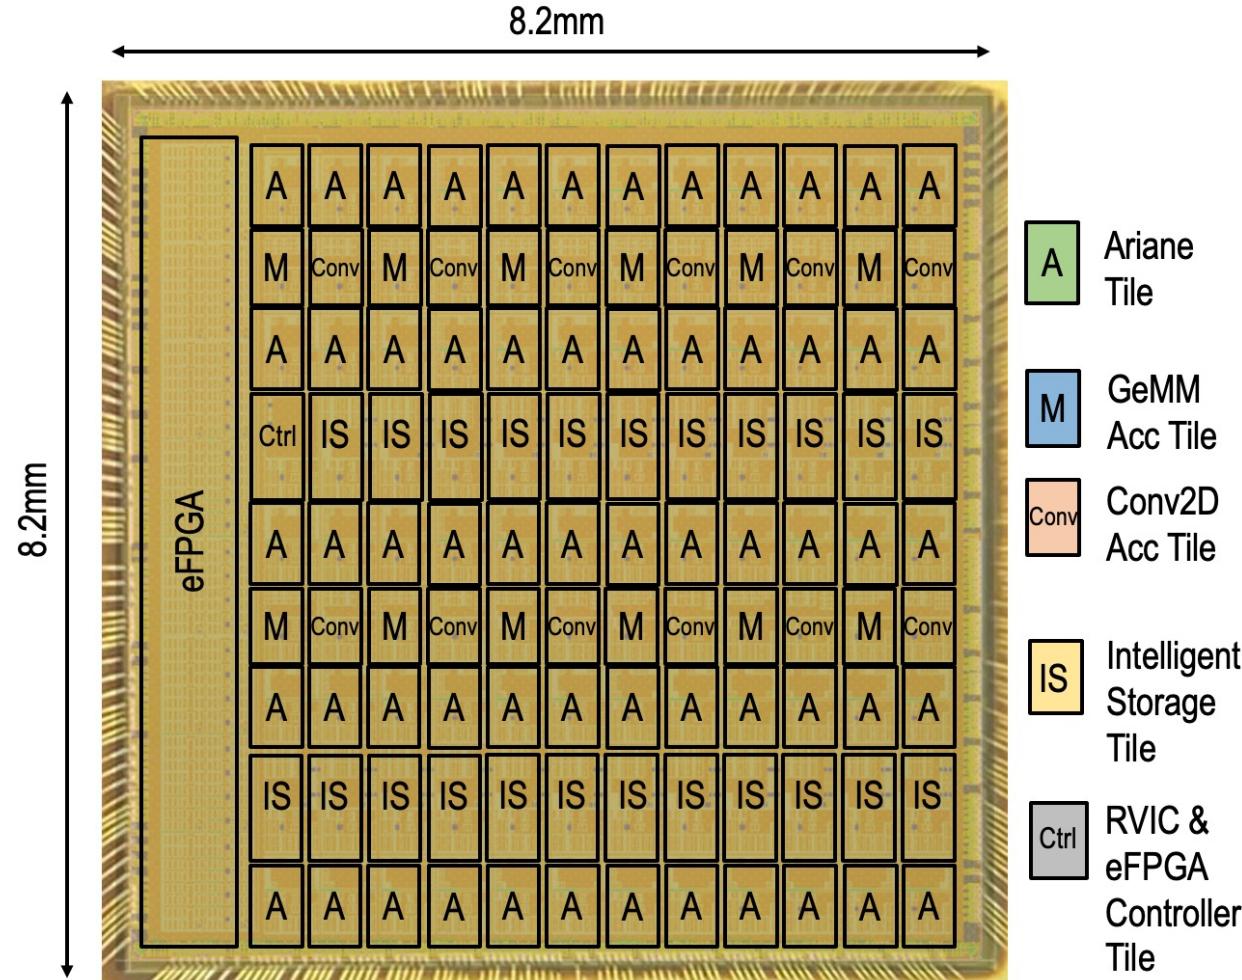
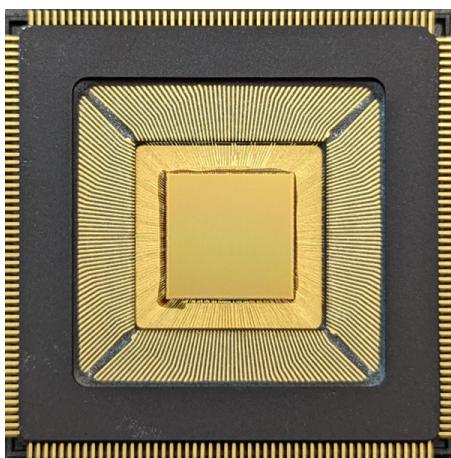
DECADES Architecture

- Accelerator-rich heterogeneous many-core architecture (108 tiles)
- 2.2B transistors, 8x8mm in 12nm process

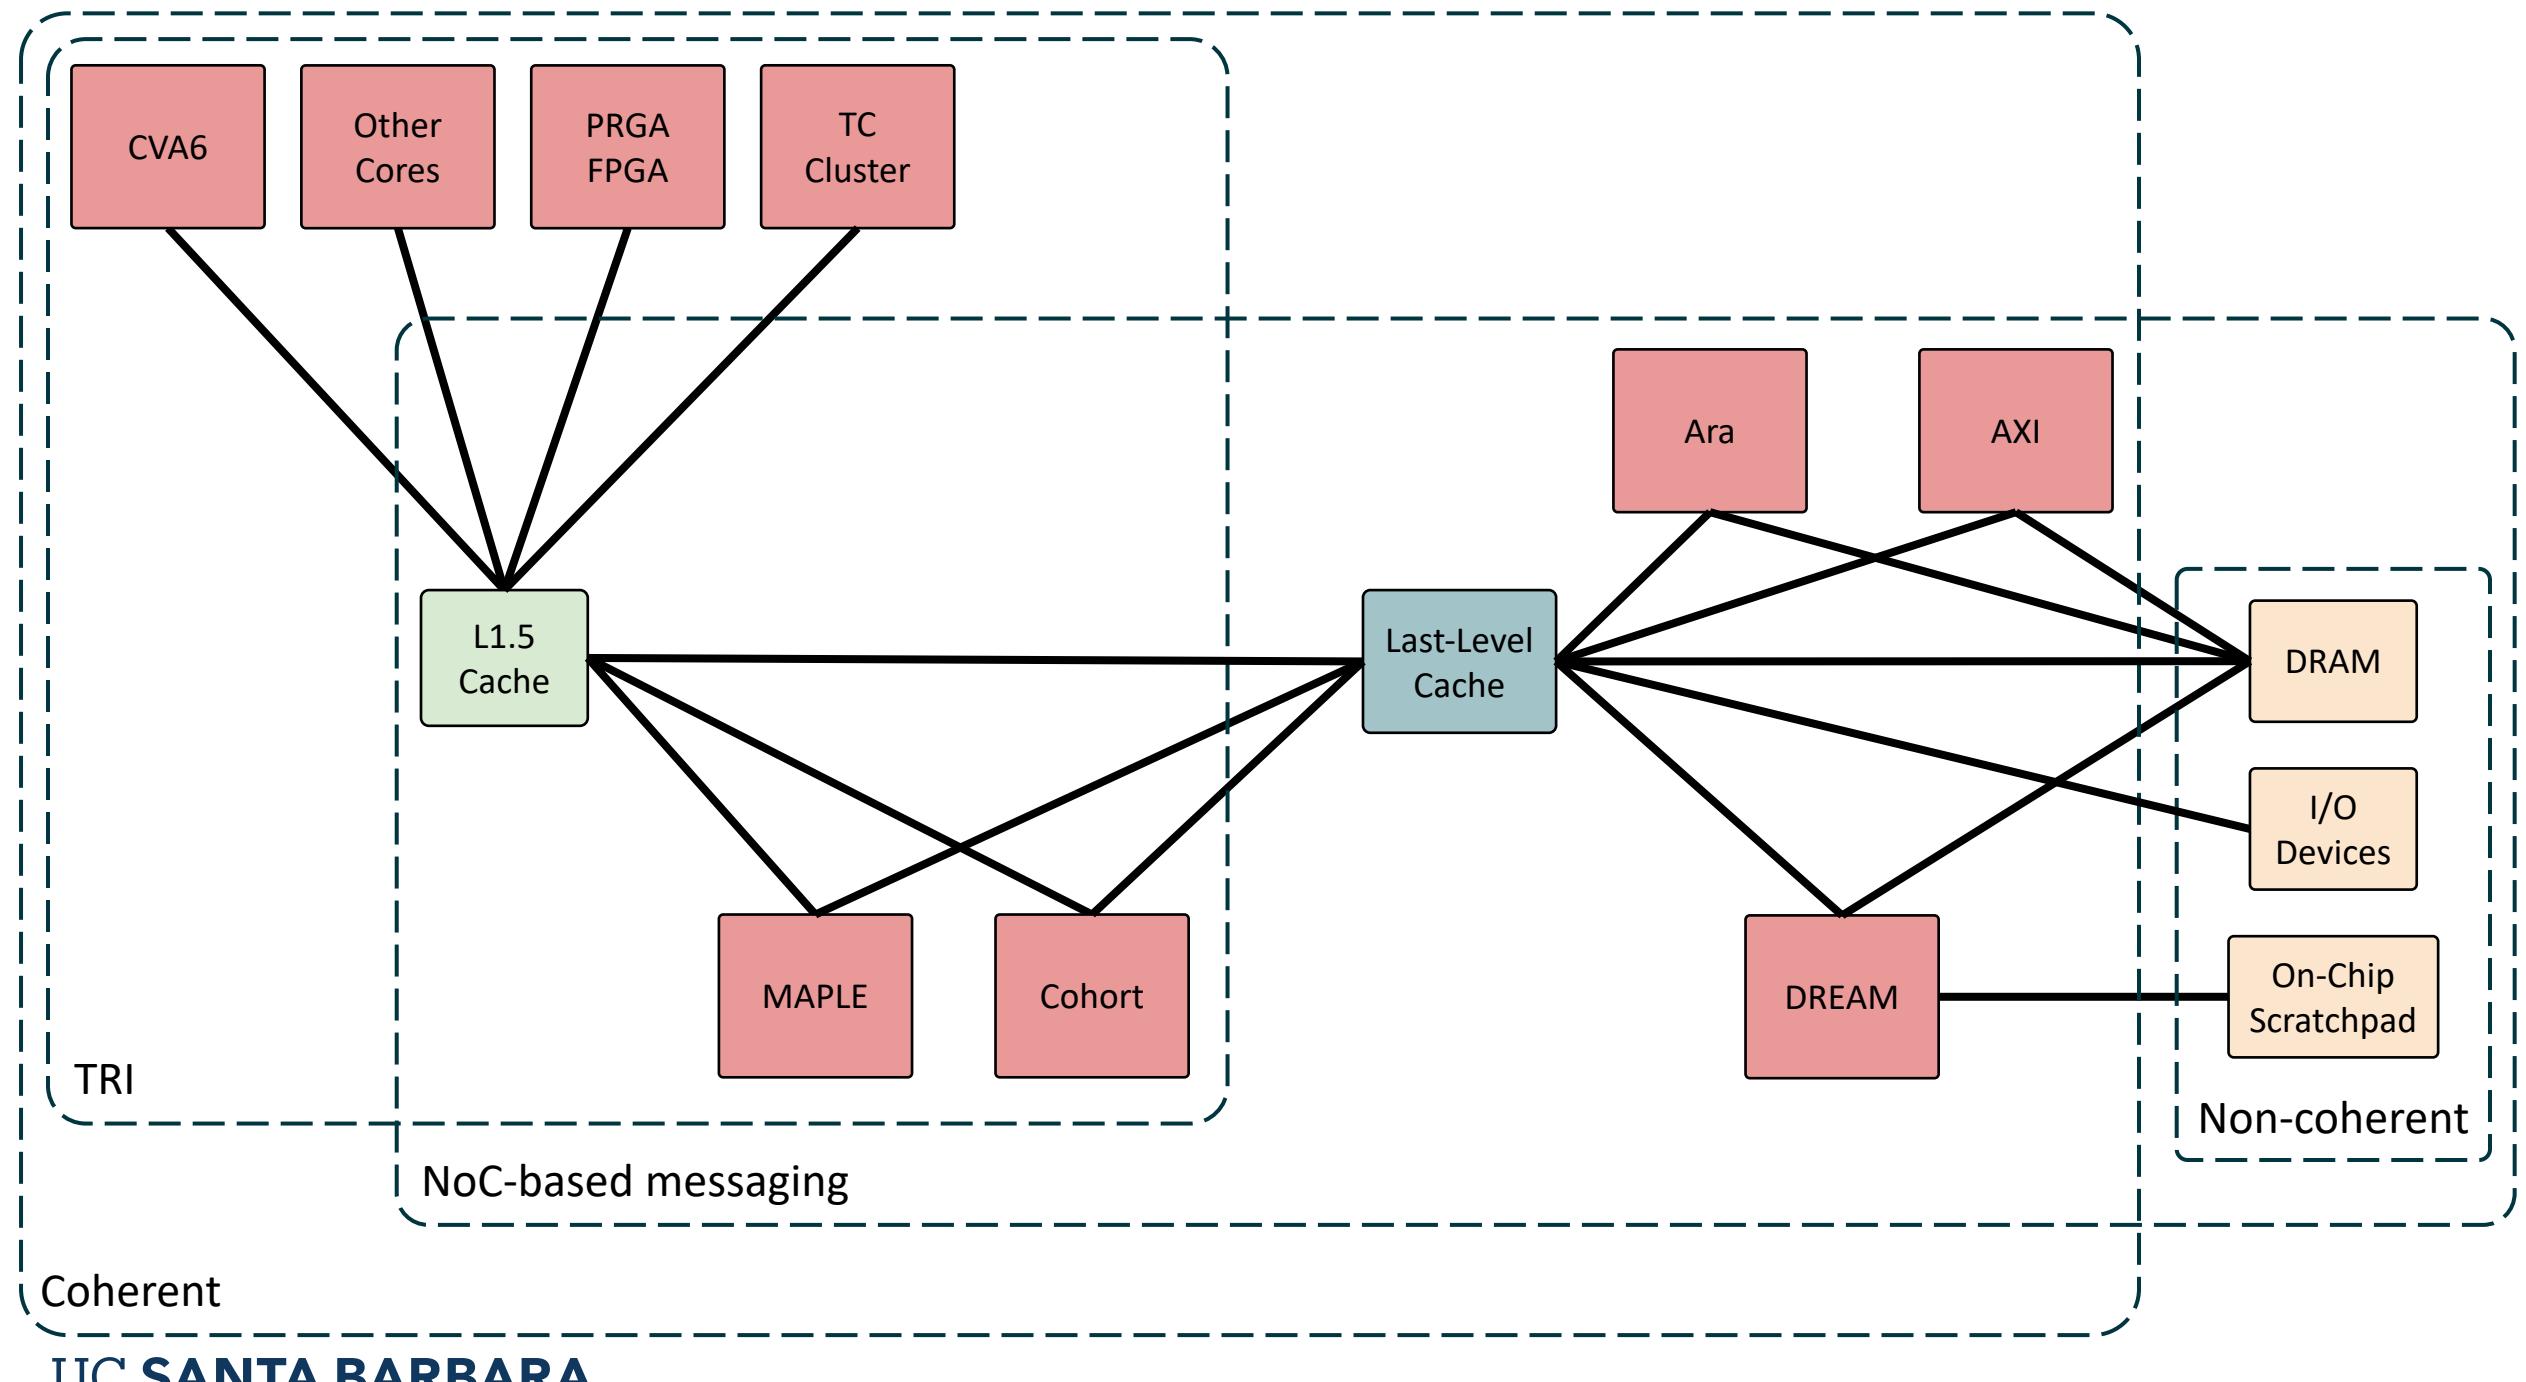
Tile type	Count
Ariane (RV64 CPU)	60
Accelerator (GeMM/Conv2D)	24
IS (Intelligent Storage)	23
Controller	1
eFPGA	1



DECADES Architecture

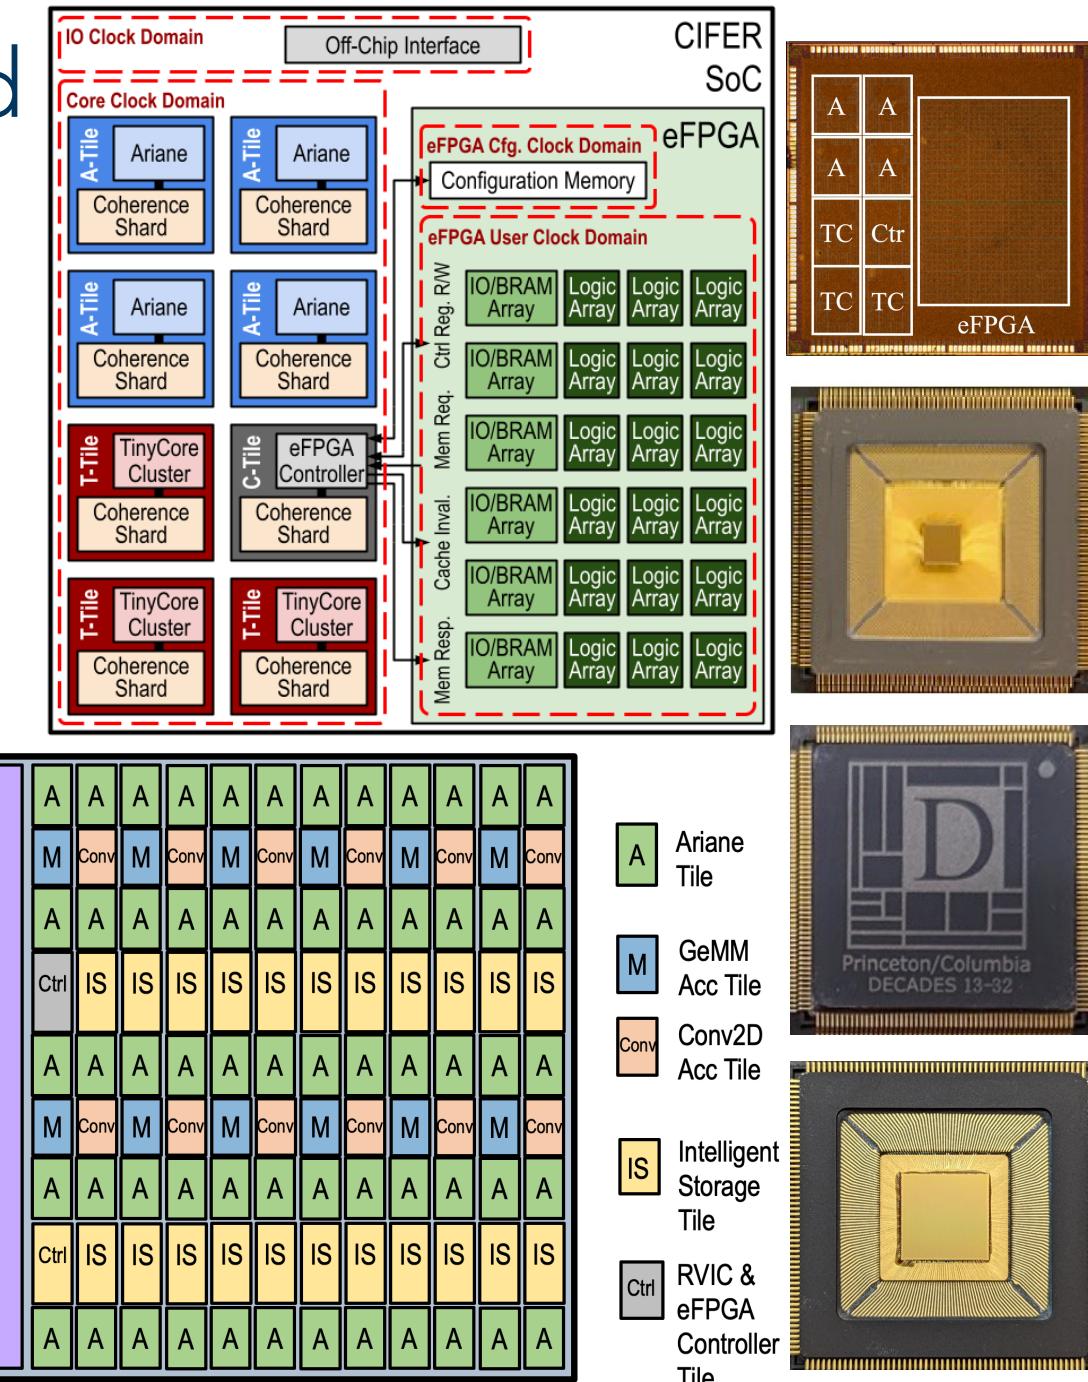


- Taped out in GF12
- 2.2 billion transistors
- All work was completed by a group of graduate students and postdoc

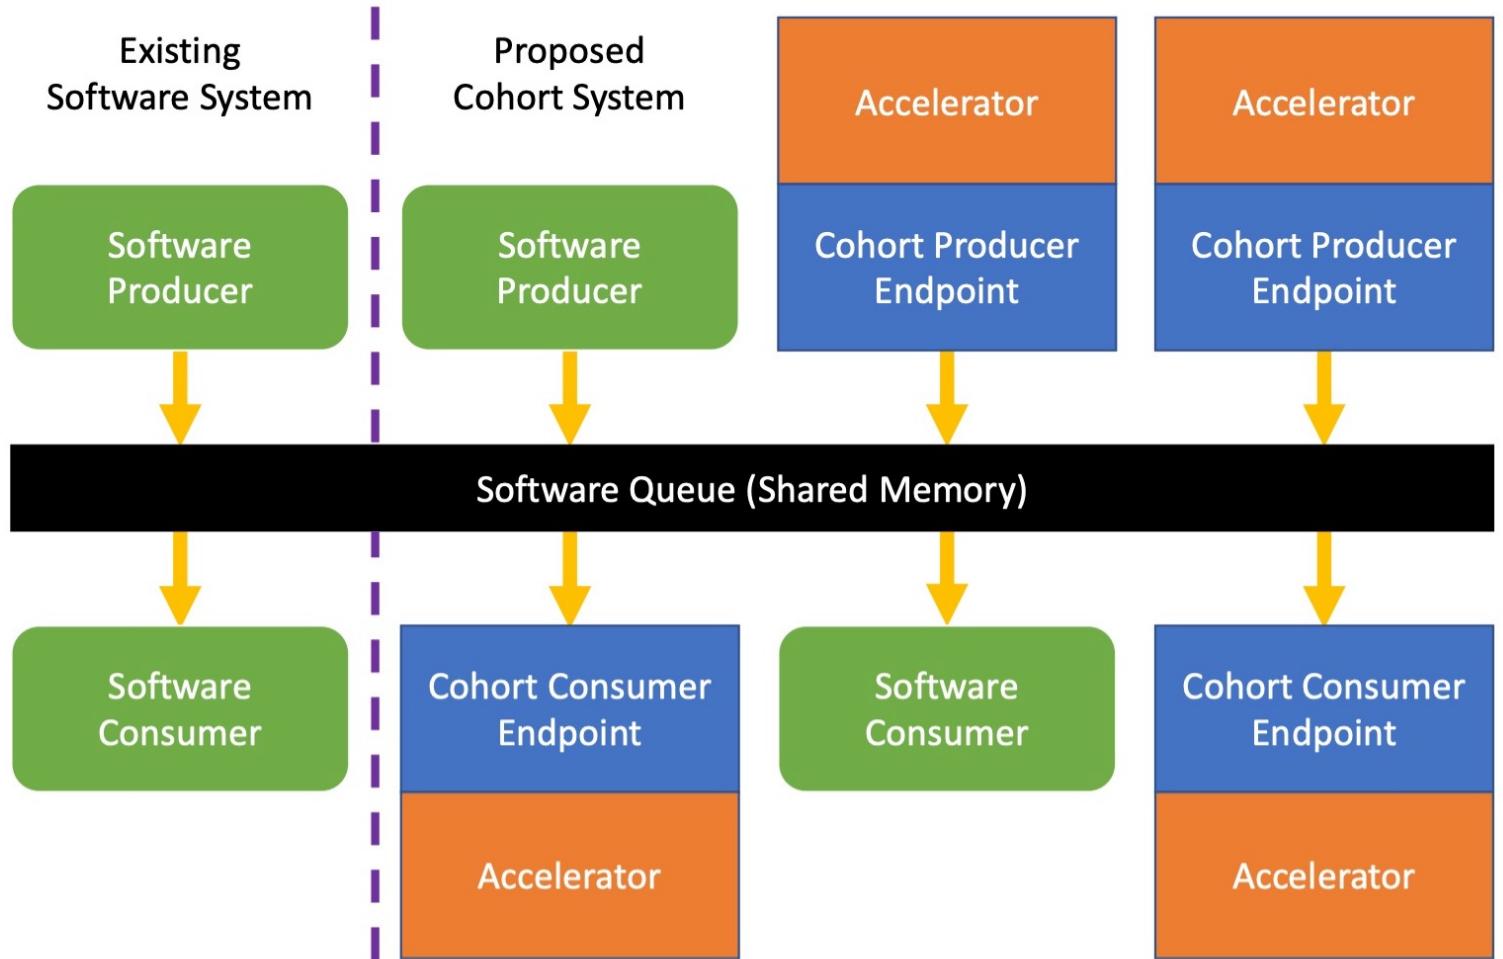


Designing modern chips is hard

- Scale and heterogeneity have become blockers
- Simulation speed plummets as design size grows
- Programming the many heterogeneous elements requires specialised code and knowledge
- What can we learn from software?

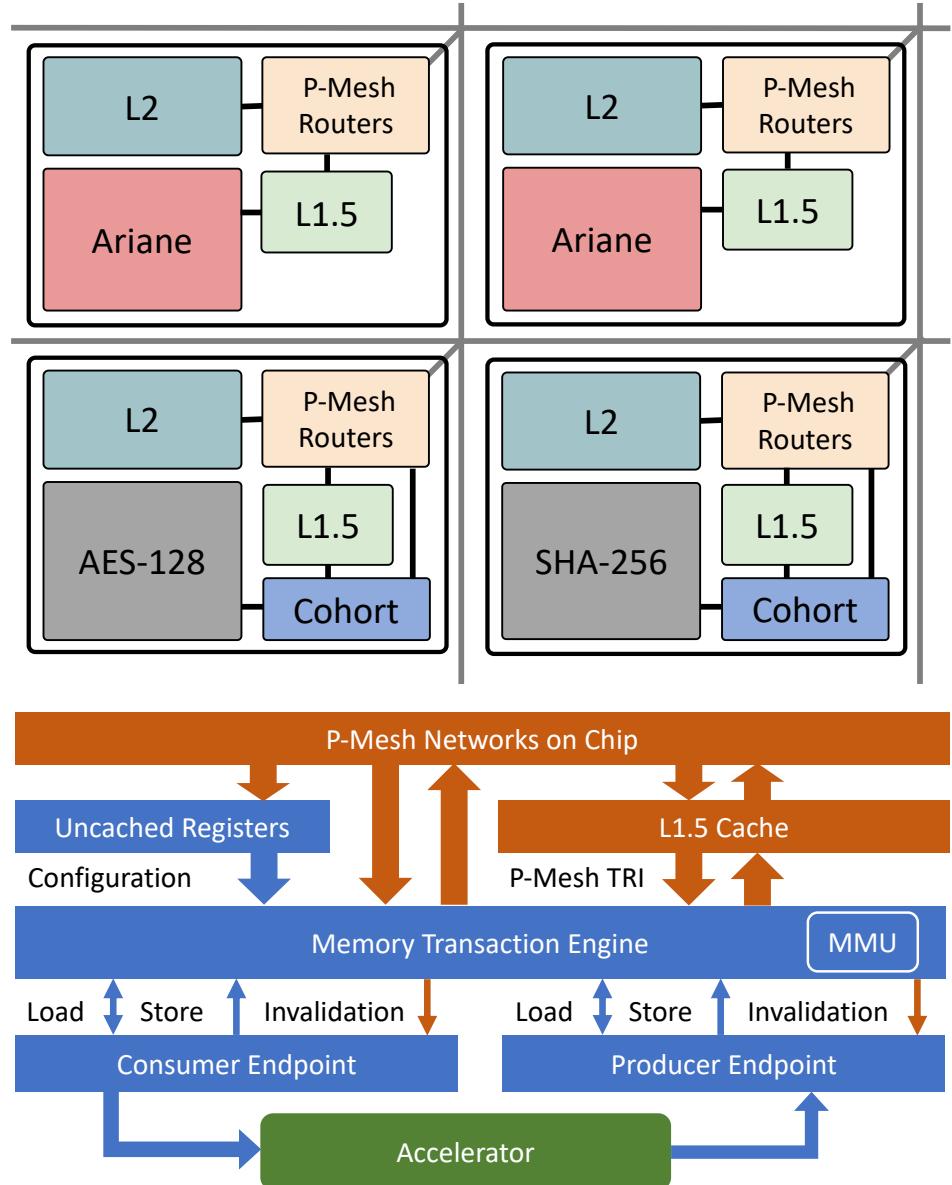


Software Oriented Acceleration



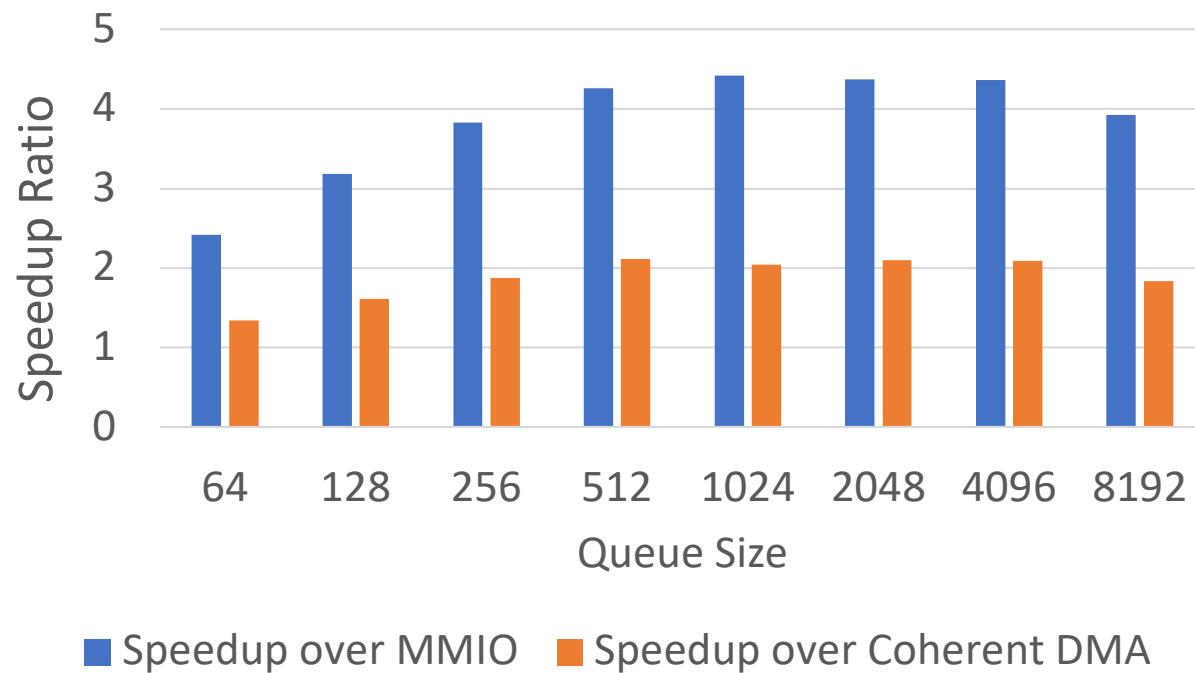
Bird's Eye View of Cohort

```
fifo_t pf = fifo_init(...);  
fifo_t cf = fifo_init(...);  
cohort_register(acc, pf, cf);  
push(data, pf);  
int result = pop(cf);
```

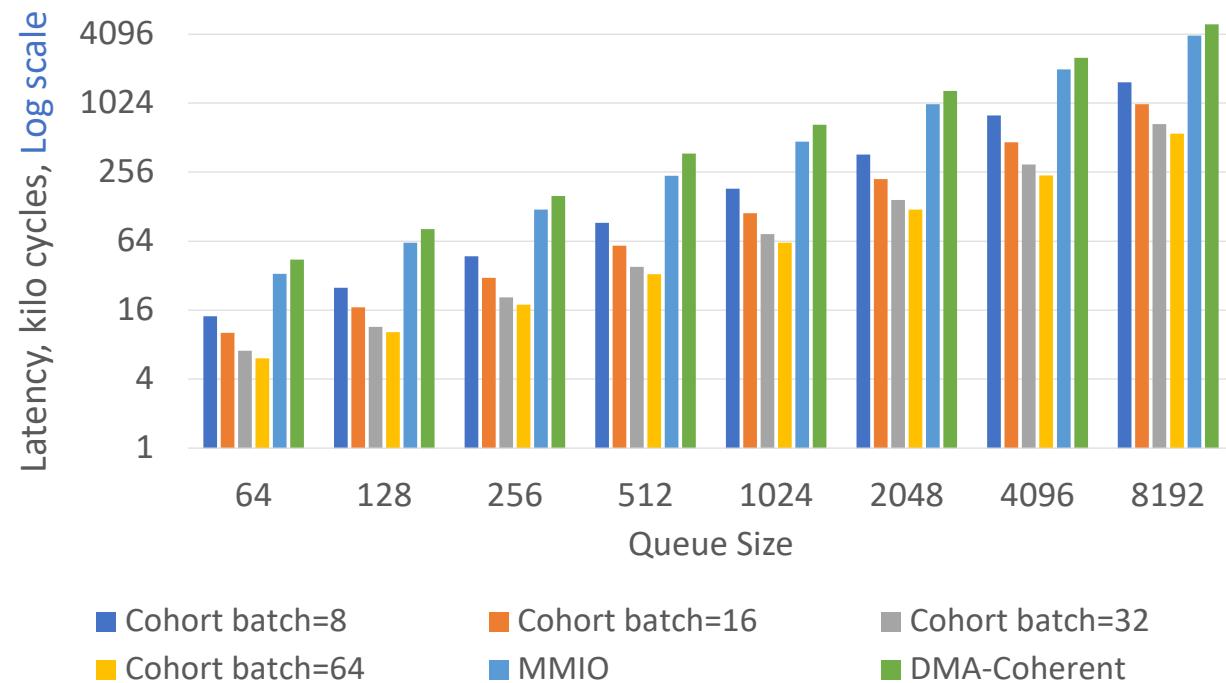


Evaluation: SHA Accelerator

IPC Performance with SHA accelerator



Program Latency with SHA accelerator



UC SANTA BARBARA